

# Kenny Sheridan



Infrastructure Product Engineer for AI, robotics, edge, agentic workflows, and supercomputing systems.

Systems engineer with 10+ years of experience turning complex infrastructure into production-ready platforms.

I build and evaluate reproducible AI-first, robotics, edge, and datacenter systems for agentic workloads, including offline AI and HPC for regulated or contested environments, bare-metal GPU orchestration, Kubernetes control planes, distributed Raft-backed reconciliation loops, high-performance networking and storage, QEMU/KVM sandboxing, observability, benchmarking, workload-to-hardware matching, and technical stack/infrastructure viability assessments.

My experience spans medium-sized companies scaling into Fortune 500 operations, as well as pre-seed through post-Series A startups valued from \$300M to \$1B. I work well in startup environments where priorities change quickly and ideas need to become shipped systems without losing reliability or technical depth. I'm a former U.S. Marine Corps meteorology instructor with hands-on HPC delivery experience across NVIDIA and AMD platforms for AI/ML, robotics, edge, simulation, and datacenter workloads.

10K+

GPUs automated across heterogeneous providers and clusters

800G+

InfiniBand and NVMe-oF fabric experience

10+ yrs

Production infrastructure and product monetization

LOCATION **Seattle, WA** EMAIL [kennethdashensheridan@gmail.com](mailto:kennethdashensheridan@gmail.com) GITHUB [github.com/kennethdsheridan](https://github.com/kennethdsheridan) FORGEJO [forgejo.kennysheridan.io](https://forgejo.kennysheridan.io)  
CODEBERG [codeberg.org/kennysheridan](https://codeberg.org/kennysheridan) LINKEDIN [linkedin.com/in/kennethdashensheridan](https://linkedin.com/in/kennethdashensheridan)

## EXPERIENCE

### Member of Technical Staff - Infrastructure Product | Andromeda

Greater Seattle Area, Remote | March 2026 - Present | Website: [andromeda.ai](https://andromeda.ai)

- Serve as a senior product-engineering contributor shaping AI-first infrastructure product direction across robotics, edge, Kubernetes, observability, platform delivery, and customer-facing workflows.
- Delivered major infrastructure products on compressed iteration cycles, translating ambiguous product and infrastructure requirements into shipped systems, automation, and operational patterns.
- Drive Kubernetes platform work beyond observability, including deployment workflows, control-plane integration, cluster operations, developer experience, and productized infrastructure paths.
- Create rich Kubernetes and infrastructure product documentation for customer-facing workflows, designed for online hosting, mobile viewing, and clear operational adoption.
- Use Nix, Kubernetes, and QEMU/KVM to build reproducible package, sandboxing, isolation, deployment, and repeatable infrastructure validation environments.
- Guide performance engineering, system modeling, cache strategy, storage evaluation, offload analysis, and workload placement so products reach market quickly without losing operational reliability.

### Senior Supercomputing Infrastructure Engineer | San Francisco Compute Company

Greater Seattle Area, Remote | September 2024 - March 2026 | Website: [sfcompute.com](https://sfcompute.com)

- Led automated bring-up for 2,000 NVIDIA H100 GPUs, moving bare metal into operational Kubernetes clusters through a single-command Rust-based deployment workflow.
- Scaled onboarding from 8 nodes to hundreds of GPU nodes within weeks, reducing product iteration time by eliminating manual provisioning across hardware, networking, and company infrastructure integration.
- Deployed distributed supercomputing infrastructure globally for GPU marketplace capacity, with emphasis on scalable utilization, reliability, and operational repeatability.
- Built and open-sourced Rust tooling for serialized infrastructure inventory plus object-storage and network throughput profiling across bare-metal AI/HPC environments.
- Designed a private Linux-side hardware discovery and lifecycle agent for bare-metal fleet introspection, host validation, and infrastructure control-plane integration.
- Optimized compute, SDN, network fabric, high-performance storage, performance testing, custom Kubernetes controllers/operators, resource management, near-metal validation, and repeatable systems setup.

### Senior AI and HPC Infrastructure Engineer | TensorWave

Greater Seattle Area, Remote | May 2024 - July 2024 | Website: [tensorwave.com](https://tensorwave.com)

- Architected on-premises AMD MI300X GPU clusters using EPYC CPUs and RDMA over Converged Ethernet on traditional TCP/IP networks.

## SKILLS

|           |                                                     |
|-----------|-----------------------------------------------------|
| CODE      | Linux, Rust, Go, Bash, NixOS                        |
| BUILD     | Nix packaging, QEMU/KVM                             |
| CONTROL   | Kubernetes operators, GitOps, Argo CD               |
| AGENTIC   | Workflows, harnesses, MCP                           |
| DATA      | Raft reconciliation, embedded databases             |
| HARDWARE  | Bare metal, Redfish, H100, H200, B200, B300, MI300X |
| GPU       | CUDA, ROCm, NCCL                                    |
| FABRIC    | RoCE, InfiniBand, NVMe-oF, SONiC, SDN               |
| OVERLAY   | Tailscale, Netmaker                                 |
| OFFLOAD   | SmartNICs, Smart Storage cards                      |
| STORAGE   | Weka, VAST, Ceph                                    |
| TELEMETRY | OpenTelemetry, Prometheus, Grafana                  |
| PERF      | FIO, iperf3, stress-ng, MPI, PyTorch                |

## STRENGTHS

### Defense AI infrastructure

Build offline AI and HPC platforms for regulated, contested, edge, simulation, and datacenter workloads.

### Reproducible systems

Make infrastructure repeatable, auditable, and safe to operate across environments.

### GPU and edge orchestration

Automate accelerator fleets across providers, clusters, and constrained networks.

### Fabric and storage paths

Design data movement, storage access, offload paths, and telemetry for reliable workloads.

### Validation-first delivery

Benchmark, stress, profile, model, and place workloads with operational evidence.

## OPEN SOURCE

- Benchmarked NVIDIA InfiniBand and AMD RoCE designs, including high-bandwidth all\_reduce testing over 800G switching infrastructure.
- Designed vendor-agnostic AI/ML infrastructure patterns capable of scaling toward hundreds of nodes while reducing accelerator lock-in.
- Created deployment documentation for GPU cluster setup, configuration, and operational handoff.

#### ○ **Senior Hardware Infrastructure Automation Engineer | ServiceNow**

Kirkland, WA | February 2022 - December 2023 | Website: [servicenow.com](https://servicenow.com)

- Engineered HPC system testing software for distributed enterprise infrastructure, including stress validation, benchmarking, and reliability assessment workflows.
- Validated infrastructure hardware for IL5, FedRAMP, and FedRAMP High environments, including Thales SafeNet security devices.
- Led migration of internal automation from Python and Bash to Go, improving efficiency across heterogeneous hardware environments.
- Built Redfish-based SKU auditing, NIC benchmarking, and GitLab CI/CD workflows for hardware-software validation.

#### ○ **Senior Cloud Hardware Performance Test Engineer | ServiceNow**

Kirkland, WA | May 2017 - February 2022 | Website: [servicenow.com](https://servicenow.com)

- Led hardware performance testing across storage, networking, BIOS, firmware, PCIe, FPGAs, SmartNICs, Smart Storage cards, NVMe, Linux filesystems, Weka, VAST, and Ceph.
- Worked with ODMs, system engineers, CTO stakeholders, and product teams to refine infrastructure roadmaps and train engineers on repeatable test methods.

#### ○ **System Administrator | NexLevel Information Technology**

Sacramento, CA | August 2015 - May 2017 | Website: [nexlevelit.com](https://nexlevelit.com)

- Provided Tier 3 Unix and Windows server support for biometric systems serving 300+ remote clients, including storage recovery, monitoring scripts, and production baseline improvements.

#### ○ **Technical Instructor of Meteorology | U.S. Marine Corps**

Quantico, VA | May 2007 - June 2015 | Website: [marines.mil](https://marines.mil)

- Administered two modular data centers, maintained METMF(R) computing infrastructure, virtualized instructional environments, and managed WAN-connected remote sensing sites.
- Produced 300+ surface observations, 100+ forecasts, and 50+ weather warnings cited in Navy and Marine Corps Achievement Medal recognition.

### SELECTED ENGINEERING WORK

#### ○ **Repeatable AI infrastructure environments**

Nix | QEMU/KVM | Secure packaging | Model-serving cache paths

Set up deterministic infrastructure environments that use Nix for secure packages, QEMU/KVM for sandboxed validation, and caching strategies to serve AI models quickly.

#### ○ **AI-first infrastructure product delivery**

Kubernetes | Product engineering | Observability | GPU platforms

Delivered customer-facing and internal infrastructure products at 9,000+ GPU scale, balancing fast iteration, operational adoption, Kubernetes platform work, and production reliability.

#### ○ **Automated GPU bring-up and onboarding**

Rust | Kubernetes | Bare metal | NVIDIA H100

Single-command workflow that deploys hardware, joins company infrastructure, configures networking, and removes manual intervention from large-scale GPU node onboarding.

#### ○ **Infrastructure inventory and throughput profiling**

Rust | Hardware inventory | Object storage | Network profiling

Built public tooling for serialized hardware reports and portable object-storage and network throughput profiling, supporting faster validation across bare-metal AI/HPC fleets.

#### ○ **Bare-metal lifecycle agent**

Rust | Linux agents | Hardware discovery | Control-plane integration

Designed private host-side agent work for hardware discovery, lifecycle state, fleet validation, and integration with infrastructure control planes.

#### ○ **Multi-node GPU cluster networking**

AMD MI300X | NVIDIA | RoCE | InfiniBand | 800G fabrics

Designed vendor-agnostic topologies across AMD and NVIDIA accelerators, including RoCE on TCP/IP networks and InfiniBand benchmarking for large-scale AI/HPC clusters.

### Forgejo portfolio

Self-hosted forge for systems architecture and private/public engineering work.

### GitHub profile

Public Rust, Nix, GPU, MCP, and developer tooling repositories.

### Neovim tooling

Lua-based personal and professional editor configuration used since 2015.

### EDUCATION

#### **Air University / Community College of the Air Force**

Keesler AFB | 62 credits in meteorology, forecasting, and atmospheric physics

### CERTIFICATIONS

#### **Introduction to Blockchain**

Amazon Web Services, 2021

#### **Basic Instructor Course**

Community College of the Air Force, 2012

#### **Meteorology and Oceanography Analyst Forecaster**

Community College of the Air Force, 2008

### RECOGNITION

Navy and Marine Corps Achievement Medal

Good Conduct Medal

Letter of Appreciation, Director of Meteorology